National University of Singapore

CS2106 Operating System

Second Half Summary Notes

Dong Shaocong A0148008J

April 17, 2018

1 I/O System

Definition 1.1. I/O Devices

- **Communication devices**: Input only (mouse, keyboard); output only (display); Input/output (network card)
- **Storage devices**: Input/output (disk, tape); Input only (CD-ROM)

Definition 1.2. Main tasks of I/O System

- Present **logical** (abstract) view of devices (**hide**: details of hardware interface and error handling)
- Facilitate **efficient** use: overlap CPU and I/O
- Support **sharing** of devices: protection when device is shared (disk), scheduling when exclusive access needed (printer)

Definition 1.3. Block-Oriented Device Interface

- **Description**: direct access, contiguous blocks, usually fixed block size
- **Operation**:
  - **Open**: verify device is ready, prepare it for access
  - **Read**: Copy a block into main memory
  - **Write**: Copy a portion of main memory to a block
  - **Close**: Release the device
  - ***Note**: these are lower level than those of the FS
- **Application**: Used by File System and Virtual Memory System; Applications typically go through the File System

Definition 1.4. Stream-Oriented Device Interface

- **Description**: character-oriented, sequential access
- **Operation**:
  - **Open**: reserve exclusive access
  - **Get**: return next character of input stream
  - **Put**: append character to output stream
  - **Close**: release exclusive access
  - ***Note**: these too are different from those of the FS but some systems try to present a uniform view of files and devices

Definition 1.5. I/O Devices - Ouput

- **Display monitors**:
  - character or graphics oriented
  - Different data rates: 25 x 80 characters vs 800 x 600 pixels (1B allows 256 colors) Refresh 30-60 times/s for video
- **Printers (ink jet, laser)**
- **Interface**:
  - **write** to controller buffer
  - **wait** for completion
  - handle **errors**

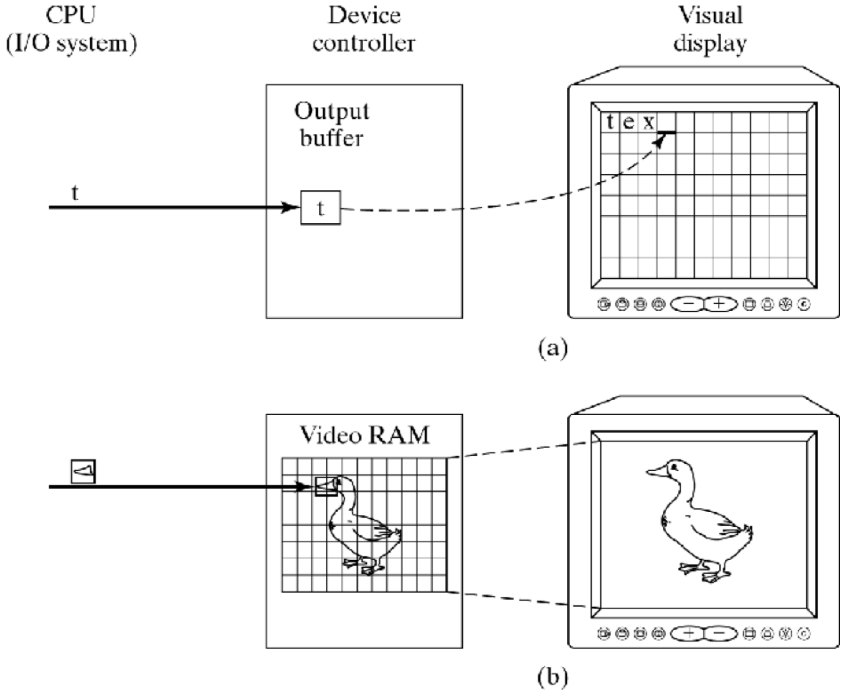

Definition 1.6. I/O Devices - Input Keyboards, pointing devices (mouse, trackball, joystick), scanners. **Interface**:

- device generates interrupt when data is ready
- read data from controller buffer
- low data rates, not time-critical

Definition 1.7. I/O Devices - Storage

- Surface, tracks/surface, sectors/track, bytes/sector
- All sectors numbered sequentially $0..(n-1)$, device controller provides mapping

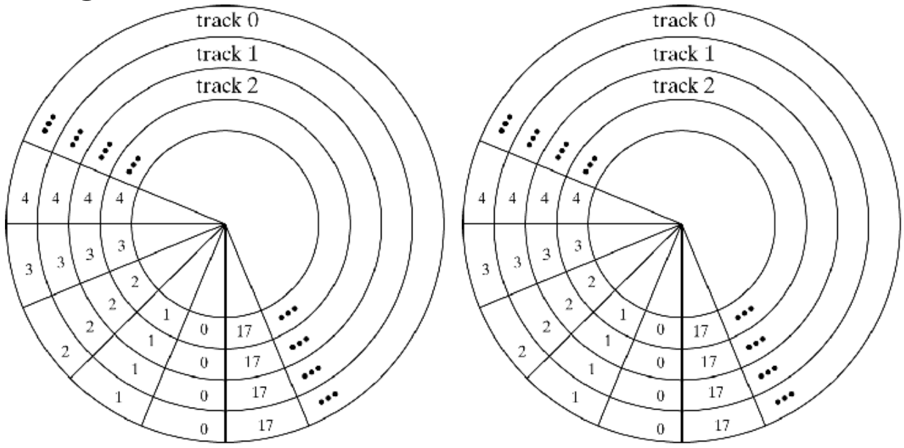

Track skew: account for seek-to-next-track to minimise rotational delay

$$\frac{\text{track to track seek time}}{\text{rotational time per track}} \times \text{sectors per track} = \text{offset}$$

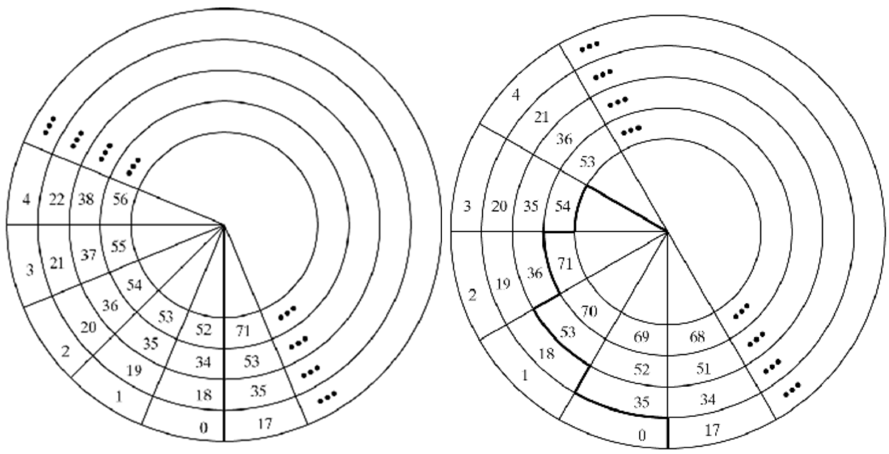

Double-sided or multiple surfaces

- Tracks with same diameter = **cylinder**
- Sectors are numbered within cylinder consecutively to **minimise seek time**

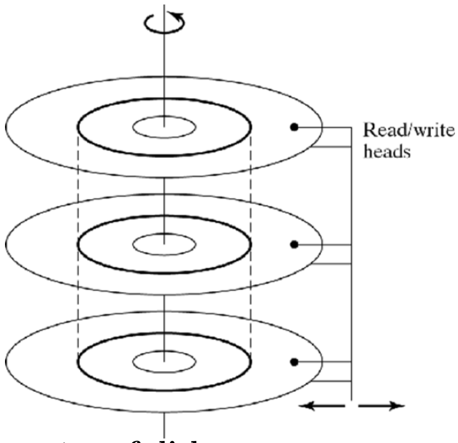

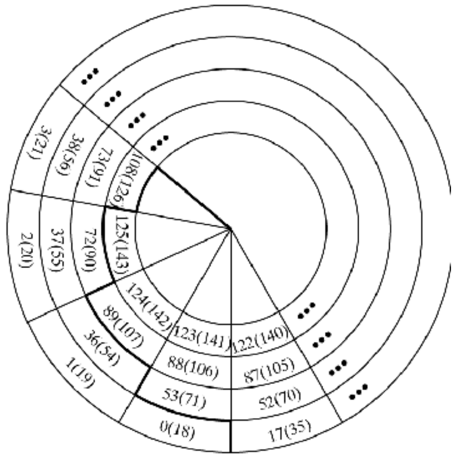

Critical issue: data transfer rates of disks

- **Sustained** rate: continuous data delivery
- **Peek** rate: : transfer once read/write head is in place; depends on rotation speed and data density

Example 1.1. Transfer rate calculation: 7200 rpm, 100 sectors/track, 512 bytes/sector

- What is the **peak** transfer rate?

$$\frac{7200}{60} \times 100 \times 512\text{byte/s}$$

- What is the **sustained** transfer rate? - Depends on file organization

Definition 1.8. I/O programming - access the I/O devices

- **Polling** (You with a broken ringer)
  - Consider a process that prints ABCDEFGH on the printer: The OS then copies character by character onto the printers latch, and the printer prints it out.
    1. Copy the first character and advance the buffers pointer.
    2. Check that the printer is ready for the next character. If not, wait. This is called “busy-waiting” or “polling”.
    3. Copy the next character. Repeat until buffer is empty.

```
copy_from_user(buffer, p, count);                   /* p is the kernel bufer */
for (i = 0; i < count; i++) {                       /* loop on every character */
    while (*printer_status_reg != READY) ;          /* loop until ready */
    *printer_data_register = p[i];                  /* output one character */
}
return_to_user( );
```

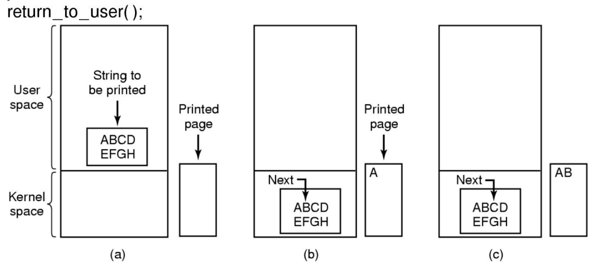

  Issue: It takes perhaps 10ms to print a character. During this time, the CPU will be busy-waiting until the printer is done printing. On a 3.2 GHz processor this is equivalent to wasting 320,000,000 instructions!

- **Interrupt-driven I/O** (You with a fully functioning phone)
  - After the string is copied, the OS will send a character to the printer, then switch to a task.
  - When the printer is done, it will interrupt the CPU by asserting one of the interrupt request (IRQ) lines on the CPU.

Main Routine

```
copy_from_user(buffer, p, count);
enable_interrupts( );
while (*printer_status_reg != READY) ;
*printer_data_register = p[0];
scheduler( );
```

(a)

Interrupt Service Routine (ISR)

```
if (count == 0) {
    unblock_user( );
} else {
    *printer_data_register = p[i];
    count = count – 1;
    i = i + 1;
}
acknowledge_interrupt( );
return_from_interrupt( );
```

(b)

- **Direct memory access** (You have an answering machine)
  - **Driver (CPU) operation to input sequence of bytes:**

```
write_reg(mm_buf, m); // give parameters
write_reg(count, n);
write_reg(opcode, read); // start op
block to wait for interrupt;
```

    - Writing opcode triggers DMA controller
    - DMA controller issues interrupt after n chars in memory
  - **Cycle Stealing**:
    - DMA controller competes with CPU for memory access
    - generally not a problem because: 1. Memory reference would have occurred anyway; 2. CPU is frequently referencing data in registers or cache, bypassing main memory.

Definition 1.9. Device Management

- **Disk Scheduling**: Requests for different blocks arrive concurrently from different processes
- Minimize **rotational delay**: re-order requests to blocks on each track to access in one rotation
- Minimize **seek time**: Conflicting goals: Minimize total travel distance; Guarantee fairness

Algorithm 1.1. Device Management

- **FIFO**: requests are processed in the order of arrival: simple, fair, but inefficient
- **SSTF** (Shortest Seek Time First): most efficient but prone to starvation
  - always go to the track thats nearest to the current positions
- **Scan** (Elevator): fair, acceptable performance
  - maintain a direction of travel
  - always proceed to the nearest track in the current direction of travel
  - if there is no request in the current direction, reverse direction

Example: assume moving from 0 to 5; then 12,4,7 arrive